

Week 04a: Search Tree Algorithms

Tree Review

1/86

Binary search trees ...

- data structures designed for $O(\log n)$ search
- consist of nodes containing item (incl. key) and two links
- can be viewed as recursive data structure (subtrees)
- have overall ordering ($\text{data}(\text{Left}) < \text{root} < \text{data}(\text{Right})$)
- insert new nodes as leaves (or as root), delete from anywhere
- have structure determined by insertion order (*worst: $O(n)$*)
- operations: insert, delete, search, ...

Balanced BSTs

Balanced Binary Search Trees

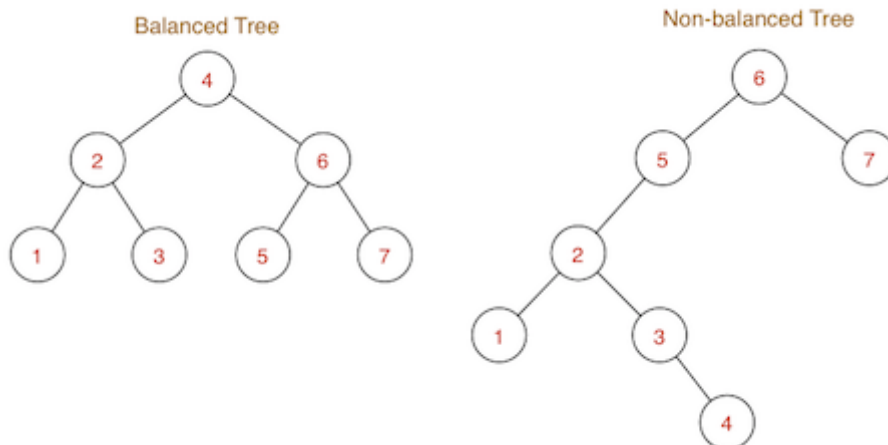
3/86

Goal: build binary search trees which have

- **minimum height** \Rightarrow minimum worst case search cost

Best balance you can achieve for tree with N nodes:

- $\text{abs}(\# \text{nodes}(\text{LeftSubtree}) - \# \text{nodes}(\text{RightSubtree})) \leq 1$, for every node
- height of $\log_2 N \Rightarrow$ worst case search $O(\log N)$



Three *strategies* to improving worst case search in BSTs:

- *randomise* — reduce chance of worst-case scenario occurring
- *amortise* — do more work at insertion to make search faster
- *optimise* — implement all operations with performance bounds

Operations for Rebalancing

4/86

To assist with rebalancing, we consider new operations:

Left rotation

- move right child to root; rearrange links to retain order

Right rotation

- move left child to root; rearrange links to retain order

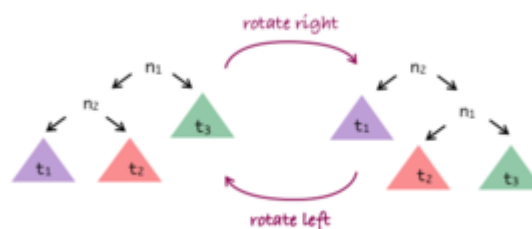
Insertion at root

- each new item is added as the new root node

Tree Rotation

5/86

In tree below: $t_1 < n_2 < t_2 < n_1 < t_3$



... Tree Rotation

6/86

Method for rotating tree T right:

- N_1 is current root; N_2 is root of N_1 's left subtree
- N_1 gets new left subtree, which is N_2 's right subtree
- N_1 becomes root of N_2 's new right subtree
- N_2 becomes new root

Left rotation: swap left/right in the above.

Cost of tree rotation: $O(1)$

... Tree Rotation

7/86

Algorithm for right rotation:

```

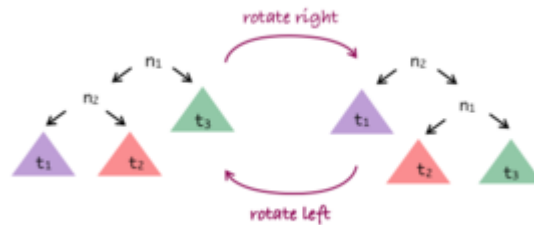
rotateRight( $n_1$ ):
|   Input   tree  $n_1$ 
|   Output  $n_1$  rotated to the right
|
|   if  $n_1$  is empty or left( $n_1$ ) is empty then
|       return  $n_1$ 
|   end if
|    $n_2 = \text{left}(n_1)$ 
|    $\text{left}(n_1) = \text{right}(n_2)$ 

```

```

|   right(n2)=n1
|   return n2

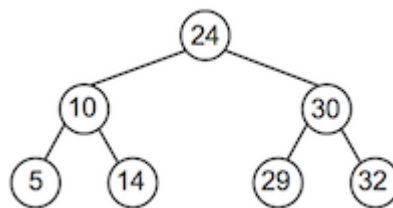
```



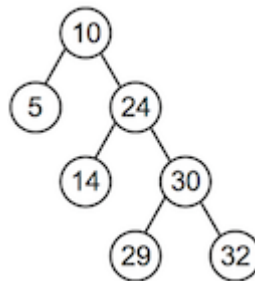
Exercise #1: Tree Rotation

8/86

Consider the tree t :



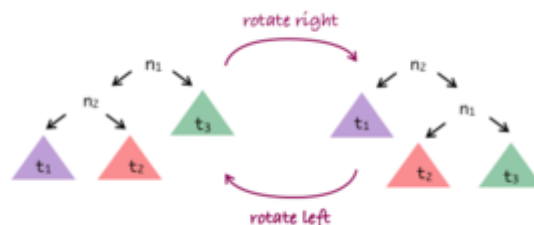
Show the result of `rotateRight(t)`



Exercise #2: Tree Rotation

10/86

Write the algorithm for left rotation



```

rotateLeft(n2):
|   Input  tree n2
|   Output n2 rotated to the left
|
|   if n2 is empty or right(n2) is empty then

```

```

|   return n2
| end if
| n1=right(n2)
| right(n2)=left(n1)
| left(n1)=n2
| return n1

```

Insertion at Root

12/86

Previous description of BSTs inserted at leaves.

Different approach: insert new item at root.

Potential disadvantages:

- large-scale rearrangement of tree for each insert

Potential advantages:

- recently-inserted items are close to root
- low cost if recent items more likely to be searched

... Insertion at Root

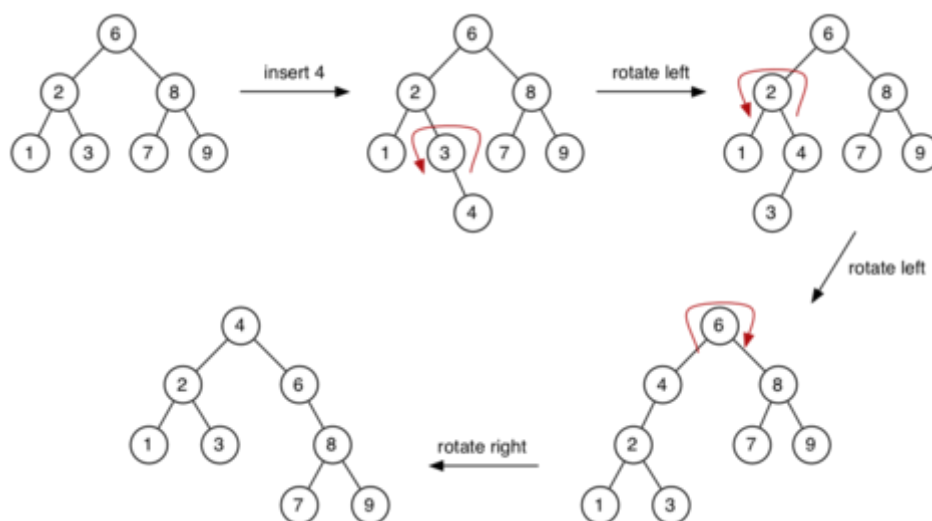
13/86

Method for inserting at root:

- base case:
 - tree is empty; make new node and make it root
- recursive case:
 - insert new node as root of appropriate subtree
 - lift new node to root by rotation

... Insertion at Root

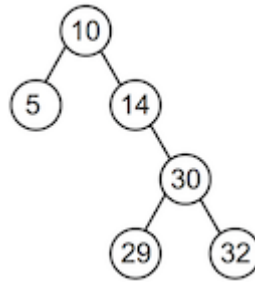
14/86



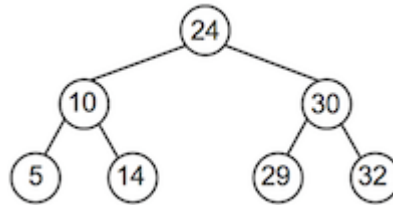
Exercise #3: Insertion at Root

15/86

Consider the tree t :



Show the result of `insertAtRoot(t , 24)`



... Insertion at Root

17/86

Analysis of insertion-at-root:

- same complexity as for insertion-at-leaf: $O(\text{height})$
- tendency to be balanced, but no balance guarantee
- benefit comes in searching
 - for some applications, search favours recently-added items
 - insertion-at-root ensures these are close to root
- could even consider "move to root when found"
 - effectively provides "self-tuning" search tree

Rebalancing Trees

18/86

An approach to balanced trees:

- insert into leaves as for simple BST
- periodically, rebalance the tree

Question: how frequently/when/how to rebalance?

`NewTreeInsert(tree, item):`

```

Input tree, item
Output tree with item randomly inserted

t=insertAtLeaf(tree,item)
if #nodes(t) mod k = 0 then
    t=rebalance(t)
end if
return t
  
```

E.g. rebalance after every 20 insertions \Rightarrow choose $k=20$

Note: To do this efficiently we would need to change tree data structure and basic operations:

```

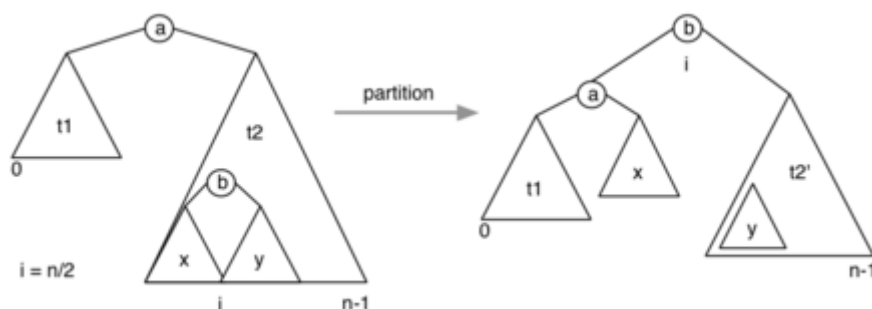
typedef struct Node {
    int data;
    int nnodes; // #nodes in my tree
    Tree left, right; // subtrees
} Node;

```

... Rebalancing Trees

19/86

How to rebalance a BST? Move median item to root.



... Rebalancing Trees

20/86

Implementation of rebalance:

```

rebalance(t):
    Input  tree t with n nodes
    Output t rebalanced

    if n ≥ 3 then
        t = partition(t, ⌊n/2⌋) // put node with median key at root
        left(t) = rebalance(left(t)) // then rebalance each subtree
        right(t) = rebalance(right(t))
    end if
    return t

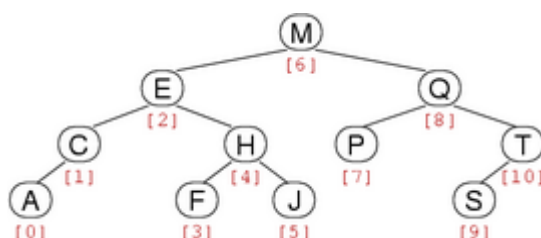
```

... Rebalancing Trees

21/86

New operation on trees:

- partition(tree, i)**: re-arrange tree so that element with index i becomes root

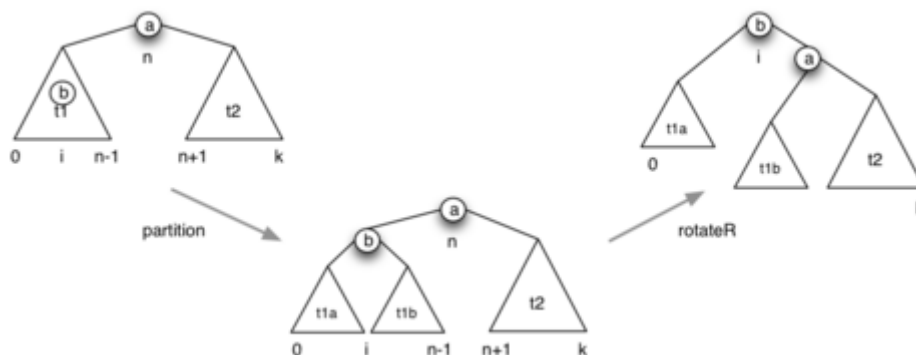


For tree with N nodes, indices are $0 \dots N-1$

... Rebalancing Trees

22/86

Partition: moves i^{th} node to root



... Rebalancing Trees

23/86

Implementation of partition operation:

```

partition(tree, i):
    Input tree with n nodes, index i
    Output tree with item #i moved to the root

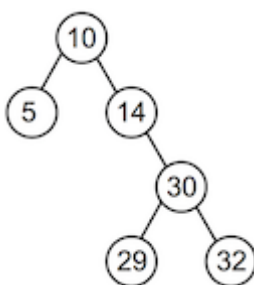
    m = #nodes(left(tree))
    if i < m then
        left(tree) = partition(left(tree), i)
        tree = rotateRight(tree)
    else if i > m then
        right(tree) = partition(right(tree), i - m - 1)
        tree = rotateLeft(tree)
    end if
    return tree
  
```

Note: $\text{size}(\text{tree}) = n$, $\text{size}(\text{left}(\text{tree})) = m$, $\text{size}(\text{right}(\text{tree})) = n - m - 1$ (why -1?)

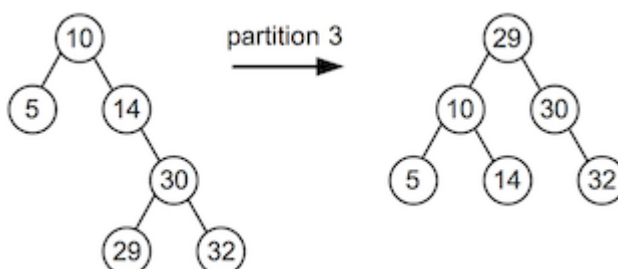
Exercise #4: Partition

24/86

Consider the tree t :



Show the result of $\text{partition}(t, 3)$



... Rebalancing Trees

Analysis of rebalancing: visits every node $\Rightarrow O(N)$

Cost means not feasible to rebalance after each insertion.

When to rebalance? ... Some possibilities:

- after every k insertions
- whenever "imbalance" exceeds threshold

Either way, we tolerate worse search performance for periods of time.

Does it solve the problem? ... Not completely \Rightarrow Solution: real balanced trees (later)

随机化的

Randomised BST Insertion

27/86

Effects of order of insertion on BST shape:

- best case (for at-leaf insertion): keys inserted in pre-order (median key first, then median of lower half, median of upper half, etc.)
- worst case: keys inserted in ascending/descending order
- average case: keys inserted in *random* order $\Rightarrow O(\log_2 n)$

Tree ADT has no control over order that keys are supplied.

Can the algorithm itself introduce some *randomness*?

In the hope that this randomness helps to balance the tree ...

... Randomised BST Insertion

28/86

How can a computer pick a number at random?

- it cannot

Software can only produce *pseudo random numbers*.

- a pseudo random number is one that is predictable
 - (although it may appear unpredictable)
- \Rightarrow implementation may deviate from expected theoretical behaviour
 - (more on this in week 10)

... Randomised BST Insertion

29/86

- Pseudo random numbers in C:

```
rand() // generates random numbers in the range 0 .. RAND_MAX
```

where the constant `RAND_MAX` is defined in `stdlib.h`
(depends on the computer: on the CSE network, `RAND_MAX` = 2147483647)

To convert the return value of `rand()` to a number between 0 .. RANGE

- compute the remainder after division by `RANGE+1`

... Randomised BST Insertion

30/86

Approach: normally do leaf insert, randomly do root insert.

```
insertRandom(tree,item)
|   Input   tree, item
|   Output tree with item randomly inserted
|
|   if tree is empty then
|       return new node containing item
|   end if
|   //  $p/q$  chance of doing root insert
|   if random number mod  $q < p$  then
|       return insertAtRoot(tree,item)
|   else
|       return insertAtLeaf(tree,item)
|   end if
```

E.g. 30% chance \Rightarrow choose $p=3, q=10$

... Randomised BST Insertion

31/86

Cost analysis:

- similar to cost for inserting keys in random order: $O(\log_2 n)$
- does not rely on keys being supplied in random order

Approach can also be applied to deletion:

- standard method promotes inorder successor to root
- for the randomised method ...
 - promote inorder successor from right subtree, OR
 - promote inorder predecessor from left subtree

Splay Trees

Splay Trees

33/86

A kind of "self-balancing" tree ...

Splay tree insertion modifies insertion-at-root method:

- by considering parent-child-grandchild (three level analysis)
- by performing double-rotations based on p-c-g orientation

The idea: appropriate double-rotations improve tree balance.

... Splay Trees

34/86

Splay tree implementations also do *rotation-in-search*:

- by performing double-rotations also when searching

The idea: provides similar effect to periodic rebalance.

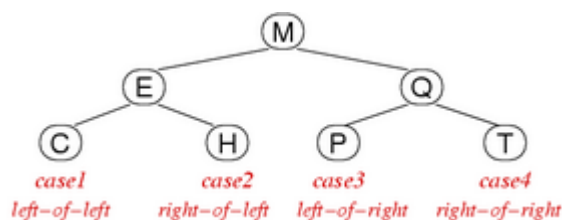
⇒ improves balance but makes search more expensive

... Splay Trees

35/86

Cases for splay tree double-rotations:

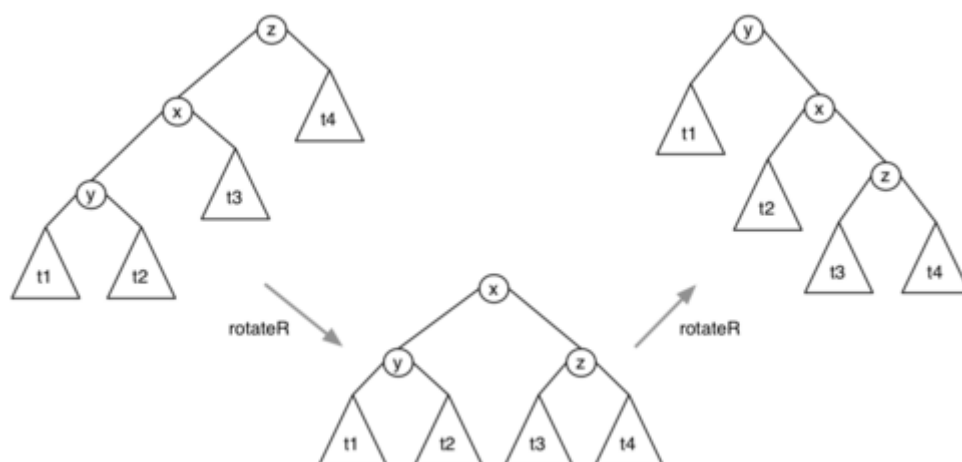
- case 1: grandchild is left-child of left-child ⇒ double right rotation from top
- case 2: grandchild is right-child of left-child
- case 3: grandchild is left-child of right-child
- case 4: grandchild is right-child of right-child ⇒ double left rotation from top



... Splay Trees

36/86

Double-rotation case for left-child of left-child ("zig-zig"):

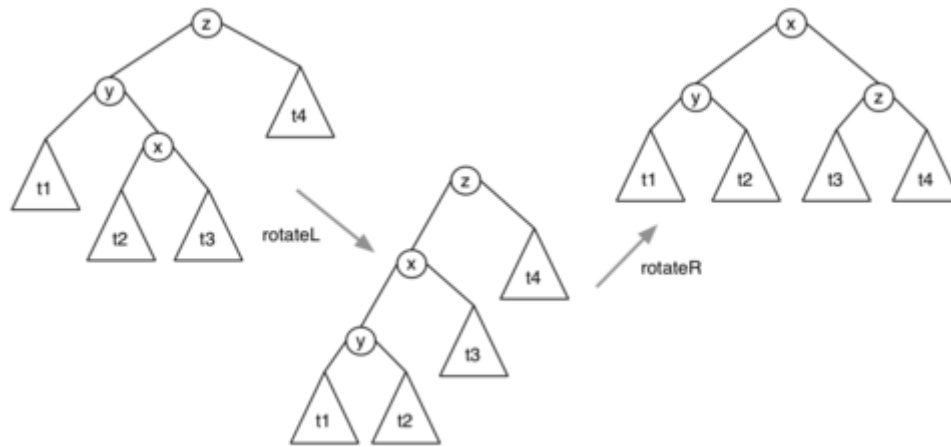


Note: both rotations at the root (unlike insertion-at-root)

... Splay Trees

37/86

Double-rotation case for right-child of left-child ("zig-zag"):



Note: rotate subtree first (like insertion-at-root)

... Splay Trees

38/86

Algorithm for splay tree insertion:

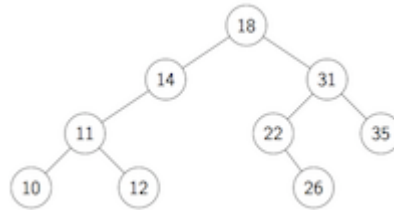
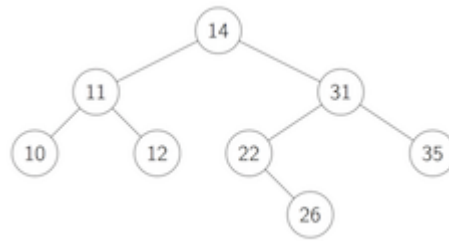
```
insertSplay(tree, item):
    Input  tree, item
    Output tree with item splay-inserted

    if tree is empty then return new node containing item
    else if item=data(tree) then return tree
    else if item<data(tree) then
        if left(tree) is empty then
            left(tree)=new node containing item
        else if item<data(left(tree)) then
            // Case 1: left-child of left-child "zig-zig"
            left(left(tree))=insertSplay(left(left(tree)), item)
            tree=rotateRight(tree)
        else if item>data(left(tree)) then
            // Case 2: right-child of left-child "zig-zag"
            right(left(tree))=insertSplay(right(left(tree)), item)
            left(tree)=rotateLeft(left(tree))
        end if
        return rotateRight(tree)
    else // item>data(tree)
        if right(tree) is empty then
            right(tree)=new node containing item
        else if item<data(right(tree)) then
            // Case 3: left-child of right-child "zag-zig"
            left(right(tree))=insertSplay(left(right(tree)), item)
            right(tree)=rotateRight(right(tree))
        else if item>data(right(tree)) then
            // Case 4: right-child of right-child "zag-zag"
            right(right(tree))=insertSplay(right(right(tree)), item)
            tree=rotateLeft(tree)
        end if
        return rotateLeft(tree)
    end if
```

Exercise #5: Splay Trees

39/86

Insert 18 into this splay tree:



... Splay Trees

41/86

Searching in splay trees:

```

searchSplay(tree,item):
  Input   tree, item
  Output address of item if found in tree
           NULL otherwise

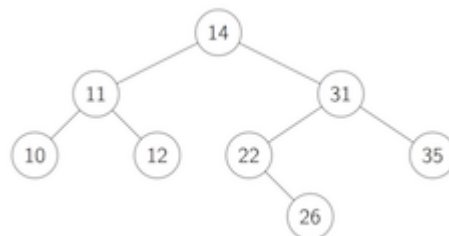
  if tree=NULL then
    return NULL
  else
    tree=splay(tree,item)
    if data(tree)=item then
      return tree
    else
      return NULL
    end if
  end if
  
```

where **splay()** is similar to **insertSplay()**,
except that it doesn't add a node ... simply moves **item** to root if found, or nearest node if not found

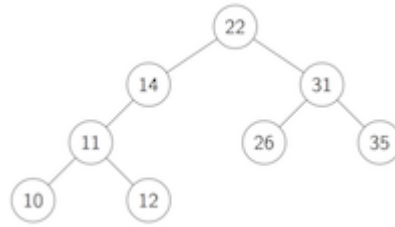
Exercise #6: Splay Trees

42/86

If we search for 22 in the splay tree



... how does this affect the tree?



... Splay Trees

44/86

Why take into account both child and grandchild?

- moves accessed node to the root
- *moves every ancestor of accessed node roughly halfway to the root*

⇒ better amortized cost than insert-at-root

... Splay Trees

45/86

Analysis of splay tree performance:

- assume that we "splay" for both insert and search
- consider: m insert+search operations, n nodes
- *Theorem.* Total number of comparisons: average $O((n+m) \cdot \log(n+m))$

Gives good overall (amortized) cost.

- insert cost not significantly different to insert-at-root
- search cost increases, but ...
 - improves balance on each search
 - moves frequently accessed nodes closer to root

But ... still has worst-case search cost $O(n)$

Real Balanced Trees

Better Balanced Binary Search Trees

47/86

So far, we have seen ...

- randomised trees ... make poor performance unlikely
- occasional rebalance ... fix balance periodically
- splay trees ... reasonable amortized performance
- but both types still have $O(n)$ worst case

Ideally, we want both average/worst case to be $O(\log n)$

- AVL trees ... fix imbalances as soon as they occur
- 2-3-4 trees ... use varying-sized nodes to assist balance
- red-black trees ... isomorphic to 2-3-4, but binary nodes

AVL Trees

AVL Trees

49/86

Invented by Georgy Adelson-Velsky and Evgenii Landis

Approach:

- insertion (at leaves) may cause imbalance
- repair balance as soon as we notice imbalance
- repairs done locally, not by overall tree restructure

A tree is unbalanced when: $\text{abs}(\text{height}(\text{left}) - \text{height}(\text{right})) > 1$

This can be repaired by at most two rotations:

- if left subtree too deep ...
 - if data inserted in left-right grandchild \Rightarrow left-rotate left subtree
 - rotate right
- if right subtree too deep ...
 - if data inserted in right-left grandchild \Rightarrow right-rotate right subtree
 - rotate left

Problem: determining height/depth of subtrees may be expensive.

... AVL Trees

50/86

Implementation of AVL insertion

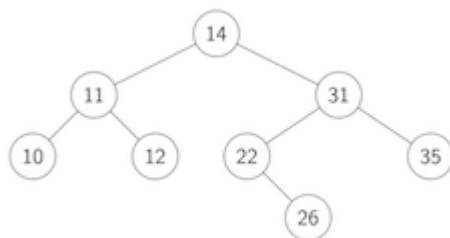
```
insertAVL(tree,item):
  Input  tree, item
  Output tree with item AVL-inserted

  if tree is empty then
    return new node containing item
  else if item=data(tree) then
    return tree
  else
    if item<data(tree) then
      left(tree)=insertAVL(left(tree),item)
    else if item>data(tree) then
      right(tree)=insertAVL(right(tree),item)
    end if
    if height(left(tree))-height(right(tree)) > 1 then
      if item>data(left(tree)) then
        left(tree)=rotateLeft(left(tree))
      end if
      tree=rotateRight(tree)
    else if height(right(tree))-height(left(tree)) > 1 then
      if item<data(right(tree)) then
        right(tree)=rotateRight(right(tree))
      end if
      tree=rotateLeft(tree)
    end if
    return tree
  end if
```

Exercise #7: AVL Trees

51/86

Insert 27 into the AVL tree



What would happen if you now insert 28?

You may like the animation at www.cs.usfca.edu/~galles/visualization/AVLtree.html

... AVL Trees

53/86

Analysis of AVL trees:

- trees are *height*-balanced; subtree depths differ by ± 1
- average/worst-case search performance of $O(\log n)$
- *require* extra data to be stored in each node ("height")
- may not be *weight*-balanced; subtree sizes may differ



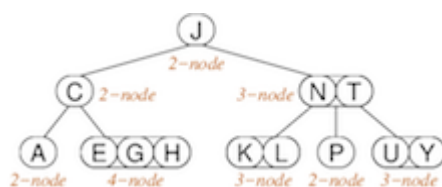
2-3-4 Trees

2-3-4 Trees

55/86

2-3-4 trees have three kinds of nodes

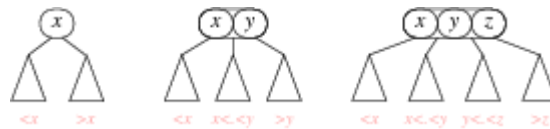
- 2-nodes, with two children (same as normal BSTs)
- 3-nodes, two values and three children
- 4-nodes, three values and four children



... 2-3-4 Trees

56/86

2-3-4 trees are ordered similarly to BSTs



In a *balanced* 2-3-4 tree:

- all leaves are at same distance from the root

2-3-4 trees grow "upwards" by splitting 4-nodes.

... 2-3-4 Trees

57/86

Possible 2-3-4 tree data structure:

```
typedef struct node {
    int      order;      // 2, 3 or 4
    int      data[3];    // items in node
    struct node *child[4]; // links to subtrees
} node;
```

... 2-3-4 Trees

58/86

Searching in 2-3-4 trees:

```
Search(tree, item):
    Input  tree, item
    Output address of item if found in 2-3-4 tree
           NULL otherwise

    if tree is empty then
        return NULL
    else
        i=0
        while i<tree.order-1 and item>tree.data[i] do
            i=i+1 // find relevant slot in data[]
        end while
        if item=tree.data[i] then // item found
            return address of tree.data[i]
        else // keep looking in relevant subtree
            return Search(tree.child[i], item)
        end if
    end if
```

... 2-3-4 Trees

59/86

2-3-4 tree searching cost analysis:

- as for other trees, worst case determined by height h
- 2-3-4 trees are always balanced \Rightarrow height is $O(\log n)$

- worst case for height: all nodes are 2-nodes
same case as for balanced BSTs, i.e. $h \approx \log_2 n$
- best case for height: all nodes are 4-nodes
balanced tree with branching factor 4, i.e. $h \approx \log_4 n$

Insertion into 2-3-4 Trees

60/86

Starting with the root node:

repeat

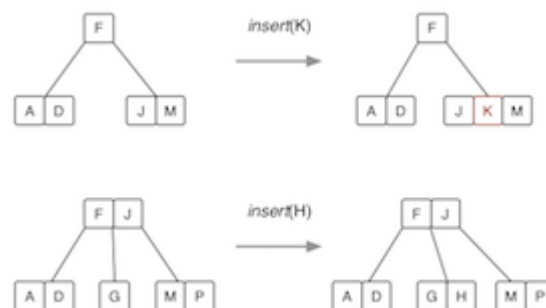
- if current node is full (i.e. contains 3 items)
 - split into two 2-nodes
 - promote middle element to parent
 - if no parent \Rightarrow middle element becomes the new root 2-node
 - go back to parent node
- if current node is a leaf
 - insert Item in this node, order++
- if current node is not a leaf
 - go to child where Item belongs

until Item inserted

... Insertion into 2-3-4 Trees

61/86

Insertion into a 2-node or 3-node:



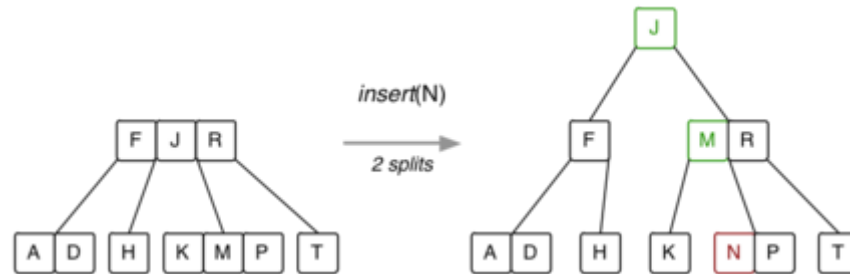
Insertion into a 4-node (requires a split):



... Insertion into 2-3-4 Trees

62/86

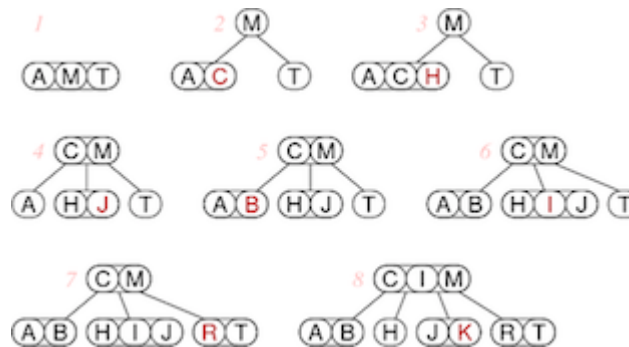
Splitting the root:



... Insertion into 2-3-4 Trees

63/86

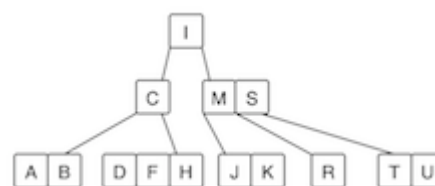
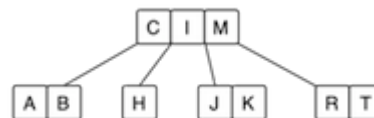
Building a 2-3-4 tree ... 7 insertions:



Exercise #8: Insertion into 2-3-4 Tree

64/86

Show what happens when D, S, F, U are inserted into this tree:



... Insertion into 2-3-4 Trees

66/86

Insertion algorithm:

```

insert(tree, item):
    Input  2-3-4 tree, item
    Output tree with item inserted

    node=root(tree), parent=NULL
    repeat
        if node.order=4 then
            promote = node.data[1]      // middle value
            nodeL   = new node containing node.data[0]
            nodeR   = new node containing node.data[2]
            if parent=NULL then
                make new 2-node root with promote, nodeL, nodeR
            else
                insert into parent
            insert into nodeL
            insert into nodeR
            return
        if item < node.data[0] then
            insert into nodeL
        else if item > node.data[2] then
            insert into nodeR
        else
            insert into node
    until not repeat
  
```

```

else
    insert promote,nodeL,nodeR into parent
    increment parent.order
end if
node=parent
end if
if node is a leaf then
    insert item into node
    increment node.order
else
    parent=node
    if item<node.data[0] then
        node=node.child[0]
    else if item<node.data[1] then
        node=node.child[1]
    else
        node=node.child[2]
    end if
end if
until item inserted

```

... Insertion into 2-3-4 Trees

67/86

Variations on 2-3-4 trees ...

Variation #1: why stop at 4? why not 2-3-4-5 trees? or M -way trees?

- allow nodes to hold up to $M-1$ items, and at least $M/2$
- if each node is a disk-page, then we have a *B-tree* (databases)
- for B-trees, depending on Item size, $M > 100/200/400$

Variation #2: don't have "variable-sized" nodes

- use standard BST nodes, augmented with one extra piece of data
- implement similar strategy as 2-3-4 trees → red-black trees.

Red-Black Trees

Red-Black Trees

69/86

Red-black trees are a representation of 2-3-4 trees using BST nodes.

- each node needs one extra value to encode link type
- but we no longer have to deal with different kinds of nodes

Link types:

- *red* links ... combine nodes to represent 3- and 4-nodes
- *black* links ... analogous to "ordinary" BST links (child links)

Advantages:

- standard BST search procedure works unmodified
- get benefits of 2-3-4 tree self-balancing (although deeper)

Red-Black Trees

70/86

Definition of a *red-black tree*

- a BST in which each node is marked red or black
- no two red nodes appear consecutively on any path
- a red node corresponds to a 2-3-4 sibling of its parent
- a black node corresponds to a 2-3-4 child of its parent

Balanced red-black tree

- all paths from root to leaf have same number of black nodes

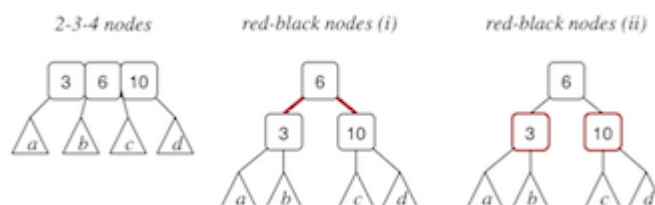
Insertion algorithm: avoids worst case $O(n)$ behaviour

Search algorithm: standard BST search

... Red-Black Trees

71/86

Representing 4-nodes in red-black trees:

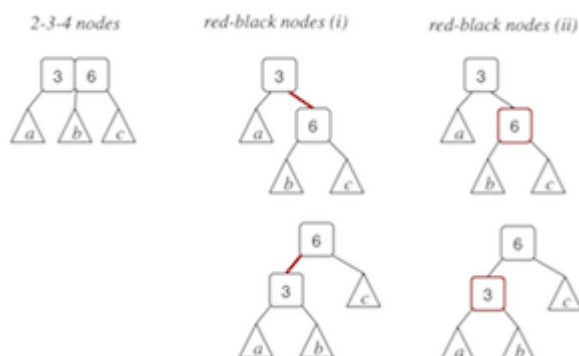


Some texts colour the links rather than the nodes.

... Red-Black Trees

72/86

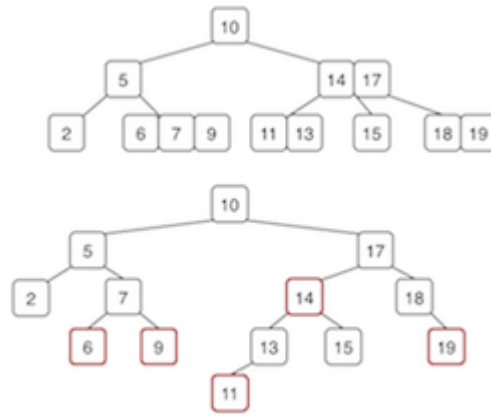
Representing 3-nodes in red-black trees (two possibilities):



... Red-Black Trees

73/86

Equivalent trees (one 2-3-4, one red-black):



... Red-Black Trees

74/86

Red-black tree implementation:

```
typedef enum {RED, BLACK} Colr;
typedef struct node *RBTree;
typedef struct node {
    int    data;    // actual data
    Colr   color;   // relationship to parent
    RBTree left;    // left subtree
    RBTree right;   // right subtree
} node;

#define color(tree) ((tree)->color)
#define isRed(tree) ((tree) != NULL && (tree)->color == RED)
```

RED = node is part of the same 2-3-4 node as its parent (sibling)

BLACK = node is a child of the 2-3-4 node containing the parent

... Red-Black Trees

75/86

New nodes are always **red**:

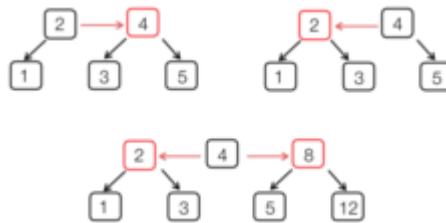
```
RBTree newNode(Item it) {
    RBTree new = malloc(sizeof(Node));
    assert(new != NULL);
    data(new) = it;
    color(new) = RED;
    left(new) = right(new) = NULL;
    return new;
}
```

... Red-Black Trees

76/86

`Node.color` allows us to distinguish links

- **black** = parent node is a "real" parent
- **red** = parent node is a 2-3-4 neighbour



... Red-Black Trees

77/86

Search method is standard BST search:

SearchRedBlack(tree, item):

```

Input tree, item
Output true if item found in red-black tree
        false otherwise

if tree is empty then
    return false
else if item < data(tree) then
    return SearchRedBlack(left(tree), item)
else if item > data(tree) then
    return SearchRedBlack(right(tree), item)
else // found
    return true
end if

```

Red-Black Tree Insertion

78/86

Insertion is more complex than for standard BSTs

- need to recall direction of last branch (L or R)
- need to recall whether parent link is red or black
- splitting/promoting implemented by rotateLeft/rotateRight
- several cases to consider depending on colour/direction combinations

... Red-Black Tree Insertion

79/86

High-level description of insertion algorithm:

insertRB(tree, item, inRight):

```

Input tree, item, inRight indicating direction of last branch
Output tree with it inserted

if tree is empty then
    return newNode(item)
else if item = data(tree) then
    return tree
end if
if left(tree) and right(tree) both are RED then
    split 4-node in a red-black tree
end if
    recursive insert a la BST, re-arrange links/colours after insert
return modified tree

```

```

insertRedBlack(tree,item):
  Input   red-black tree, item
  Output tree with item inserted

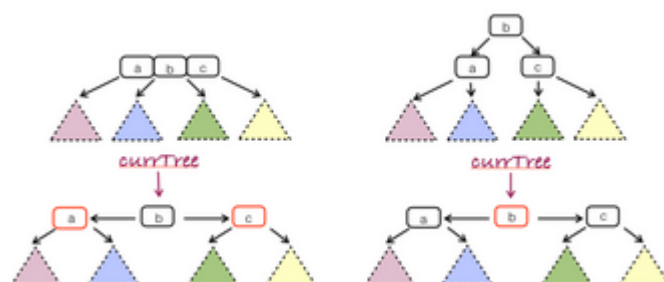
  tree=insertRB(tree,item,false)
  color(tree)=BLACK
  return tree

```

... Red-Black Tree Insertion

80/86

Splitting a 4-node, in a red-black tree:



Algorithm:

```

if isRed(left(currentTree)) and isRed(right(currentTree)) then
  color(currentTree)=RED
  color(left(currentTree))=BLACK
  color(right(currentTree))=BLACK
end if

```

... Red-Black Tree Insertion

81/86

Simple recursive insert (a la BST):



Algorithm:

```

if item<data(tree) then
  left(tree)=insertRB(left(tree),item,false)
  re-arrange links/colours after insert
else // item larger than data in root
  right(tree)=insertRB(right(tree),item,true)
  re-arrange links/colours after insert
end if

```

Not affected by colour of tree node.

... Red-Black Tree Insertion

82/86

Re-arrange links/colours after insert:

Step 1 — "normalise" direction of successive red links



Algorithm:

```

if inRight and both currentTree and left(currentTree) are red then
    currentTree=rotateRight(currentTree)
end if
  
```

Symmetrically,

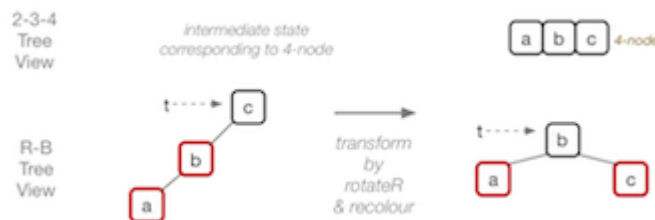
- if not inRight and both currentTree and right(currentTree) are red
⇒ left rotate currentTree

... Red-Black Tree Insertion

83/86

Re-arrange links/colours after insert:

Step 2 — two successive red links = newly-created 4-node



Algorithm:

```

if left(currentTree) and left(left(currentTree)) are red then
    currentTree=rotateRight(currentTree)
    color(currentTree)=BLACK
    color(right(currentTree))=RED
end if
  
```

Symmetrically,

- if both right(currentTree) and right(right(currentTree)) are red
⇒ left rotate currentTree, then re-colour currentTree and left(currentTree)

... Red-Black Tree Insertion

84/86

Example of insertion, starting from empty tree:

22, 12, 8, 15, 11, 19, 43, 44, 45, 42, 41, 40, 39



Red-black Tree Performance

85/86

Cost analysis for red-black trees:

- tree is well-balanced; worst case search is $O(\log_2 n)$
- insertion affects nodes down one path; max #rotations is $2 \cdot h$ (where h is the height of the tree)

Only disadvantage is complexity of insertion/deletion code.

Note: red-black trees were popularised by Sedgewick.

Summary

86/86

- Tree operations
 - tree rotation
 - tree partition
 - joining trees
- Randomised insertion
- Self-adjusting trees
 - Splay trees
 - AVL trees
 - 2-3-4 trees
 - Red-black trees
- Suggested reading:
 - Sedgewick, Ch. 12.8-12.9
 - Sedgewick, Ch. 13.1-13.4

Produced: 6 Jan 2020